

David Andrews

(908) 723-5723 | dha@xoba.com | broyojo.com | [in @broyojo](https://www.linkedin.com/in/broyojo) | [@Broyojo](https://www.youtube.com/channel/UCBroyojo)

EDUCATION

Georgia Institute of Technology, Atlanta, GA: *B.S. in Computer Science*

Aug. 2023 – May 2026

GPA: 4.0/4.0

Courses: DL, ML, Systems for ML, Intro to AI, Perception and Robotics, Algorithms, OS

EXPERIENCE

Georgia Tech, Chao Zhang Lab: *Undergraduate Researcher, Artificial Intelligence*

Aug. 2024 – Present

- Built PRM-guided tree-search reasoning systems using SGLang and vLLM caching to generate synthetic trajectories with error correction, running training runs of thousands of problems across 8-GPU machines.
- Trained multi-turn LLM agents with GRPO on SWE-bench, SWE-Smith, and SWE-Gym, adding reward-hacking defenses, tool-use checks, and rollout instrumentation to improve training stability.
- Scaled asynchronous PPO/GRPO/RLVR training using AReaL on Ray with a Kubernetes-managed sandbox server, generalizing the system to Docker/MCP tool environments for broader agent evaluation.

Georgia Tech, Qirun Zhang Lab: *Undergraduate Researcher, Systems*

Jan. 2026 – Present

- Building an LLM agent harness for automated bug finding and fixing in LLVM, with specialized tool skills for navigating large-scale systems codebases with long build times.
- Investigating agent-based commit bisection to accelerate identification of bug-introducing changes, reducing manual debugging effort in complex compiler histories.

Georgia Tech, EIC Lab: *Undergraduate Researcher, Hardware Design*

Feb. 2026 – Present

- Curating a ~250-task SWE-bench-style benchmark for Verilog/RTL hardware design derived from real-world GitHub PRs, using Verilator, Yosys, and Icarus Verilog for functional verification.
- Developing an agent harness to evaluate frontier models on industry-relevant hardware design tasks, targeting state-of-the-art performance on the benchmark.

DuckAI: *AI Researcher & Engineer*

Aug. 2023 – Present

- Engineered **DuckTrack**, a cross-platform computer-use capture tool adopted by OSWorld and OpenCUA for benchmarking and foundation model training.
- Designed and implemented **PRMBench**'s full-stack annotation platform (React, FastAPI, SQLite), improving labeling speed and inter-annotator agreement for process reward model datasets.
- Built large-scale TTS data pipelines (YouTube and podcast ingestion, upscaling, VAD, diarization) and contributed optimizations to insanely-fast-whisper to accelerate transcription for speech datasets.

SELECTED PROJECTS

Tool-Integrated RL Agent Training System

May 2025 – Present

- Engineered a general tool-using agent framework with Docker sandboxes and an MCP interface (browser, terminal, code execution), supporting multi-turn rollouts across math, search, SWE-bench, and puzzle tasks.
- Integrated AReaL for asynchronous PPO/GRPO/RLVR training to maximize throughput under high tool latency, coordinating rollout workers with sandboxes hosted on a Kubernetes cluster.

DuckTrack

Nov. 2023 – Jan. 2024

- Extended and maintained precise cross-platform capture of screen, keyboard, and mouse events, synchronizing native input APIs with OBS to achieve millisecond-level sub-frame alignment for computer-use datasets.

Emergent Misalignment

Oct. 2025 – Jan. 2026

- Replicated and investigated emergent misalignment arising from benign fine-tuning of LLMs; findings published on LessWrong.

AlphaEvolve Extensions

May 2025 – Sept. 2025

- Extended OpenEvolve, an open-source replication of Google DeepMind's AlphaEvolve, with new combinatorics domains including circle-packing and sphere-packing; built a parallelized parameter sweep system using Optuna.
- Co-contributed to an upstreamed PR adding MLIR attention kernel optimization tasks, achieving a 1.32x+ kernel speedup via evolved transformation parameters.

SKILLS

Languages: Python, Rust, C, C++, Go, JavaScript, Bash, Verilog

ML: PyTorch, vLLM, SGLang, Ray, AReaL, veRL, HuggingFace

Systems: Linux, Docker, Kubernetes, SLURM, AWS, FastAPI, SkyPilot, uv, Git